

# Project 2 Healthcare

October 4, 2021

Project 2 :Healthcare      Name: Niyojita Arun Raje      Cohort: 1 (DEC 2020)

WEEK 1: Data Exploration

```
[1]: import numpy as np,pandas as pd,matplotlib.pyplot as plt,seaborn as snss
```

```
[2]: data=pd.read_csv('G:/Simplilearn/Capstone Project/Project 2/Project 2/
↳Healthcare - Diabetes/health care diabetes.csv')
```

```
[3]: data.head()
```

```
[3]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
[4]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 768 entries, 0 to 767
```

```
Data columns (total 9 columns):
```

#	Column	Non-Null Count	Dtype
0	Pregnancies	768 non-null	int64
1	Glucose	768 non-null	int64
2	BloodPressure	768 non-null	int64
3	SkinThickness	768 non-null	int64
4	Insulin	768 non-null	int64
5	BMI	768 non-null	float64
6	DiabetesPedigreeFunction	768 non-null	float64

```

7   Age                768 non-null    int64
8   Outcome            768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB

```

```
[5]: data.describe()
```

```
[5]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin \
count	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479
std	3.369578	31.972618	19.355807	15.952218	115.244002
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000
75%	6.000000	140.250000	80.000000	32.000000	127.250000
max	17.000000	199.000000	122.000000	99.000000	846.000000

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000		768.000000	768.000000
mean	31.992578		0.471876	33.240885
std	7.884160		0.331329	11.760232
min	0.000000		0.078000	21.000000
25%	27.300000		0.243750	24.000000
50%	32.000000		0.372500	29.000000
75%	36.600000		0.626250	41.000000
max	67.100000		2.420000	81.000000

```
[6]: data.isnull().sum()
```

```
[6]: Pregnancies      0
      Glucose          0
      BloodPressure    0
      SkinThickness    0
      Insulin          0
      BMI              0
      DiabetesPedigreeFunction  0
      Age              0
      Outcome          0
      dtype: int64
```

```
[7]: data['Glucose'].values==0
```


```
[7]: array([False, False, False, False, False, False, False, False, False,  
          False, False, False, False, False, False, False, False, False,  
          False, False, False, False, False, False, False, False, False,  
          False, False, False, False, False, False, False, False, False,  
          False, False, False, False, False, False, False, False, False,
```

[illegible]

[illegible]

```
[8]: plt.xlabel('Glucose Class')
      data['Glucose'].plot.hist()
      print("Datatype of Glucose is:",data['Glucose'].dtypes)
```

Datatype of Glucose is: int64



We can see that there are 0 value data and Glucose cannot be 0. Hence replacing 0 with mean of Glucose class

```
[9]: data['Glucose']=data['Glucose'].replace(0,data['Glucose'].mean())
```

```
[10]: data['BloodPressure'].values==0
```


[illegible]

[illegible]

```
False, False, False, False, True, False, False, False, False,
False, False, False, False, False, False, False, True, False,
False, True, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, True, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, True, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, True, False, False, False, False,
False, True, False, False, True, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False]
```

```
[11]: plt.xlabel('Blood Pressure Class')
      data['BloodPressure'].plot.hist()
      print("Datatype of BloodPressure is:", data['BloodPressure'].dtypes)
```

Datatype of BloodPressure is: int64



We can see that there are 0 value data and BloodPressure cannot be 0.Hence replacing 0 with mean of BloodPressure class

```
[12]: data['BloodPressure']=data['BloodPressure'].replace(0,data['BloodPressure'].  
      ↪mean())
```

```
[13]: data['SkinThickness'].values==0
```

```
[13]: array([False, False,  True, False, False,  True, False,  True, False,  
          True,  True,  True,  True, False, False,  True, False,  True,  
          False, False, False,  True,  True, False, False, False,  True,  
          False, False,  True, False, False, False,  True, False, False,  
          True, False, False, False, False,  True, False, False,  True,  
          False,  True, False, False,  True, False, False, False, False,  
          False, False, False, False,  True, False,  True,  True,  True,  
          False,  True, False, False,  True, False, False, False, False,  
          True, False, False,  True, False, False, False, False, False,  
          True, False, False,  True, False, False, False, False, False,  
          False,  True,  True,  True, False,  True, False,  True, False,  
          False, False, False, False, False,  True, False,  True,  True,  
          True, False, False, False, False, False,  True,  True, False,  
          False, False, False,  True, False,  True, False, False, False,  
          False, False, False,  True, False,  True, False, False,  True,  
          False, False, False,  True, False,  True, False, False,  True,  
          False, False, False, False,  True, False, False,  True, False,  
          False,  True, False, False, False, False, False, False, False,  
          False, False,  True, False, False,  True,  True, False,  True,  
          False, False, False, False, False,  True,  True, False,  True,  
          False, False, False, False, False,  True, False, False,  True,  
          False, False, False,  True, False, False, False,  True,  True,  
          False, False, False,  True, False, False, False,  True,  True,  
          True, False, False,  True,  True, False, False, False, False,  
          False,  True, False,  True,  True, False, False, False,  True,  
          False,  True, False, False,  True,  True, False, False, False,  
          False, False, False, False, False, False,  True, False, False,  
          False, False,  True,  True, False, False,  True,  True, False,  
          False, False, False, False, False, False, False, False, False,  
          False, False,  True, False,  True, False, False, False, False,  
          False, False, False,  True, False, False, False, False,  True,
```






```
False, True, True, True, False, False, True, False, False,
False, True, False])
```

```
[14]: plt.xlabel('SkinThickness')
data['SkinThickness'].plot.hist()
print("Datatype of SkinThickness is:",data['SkinThickness'].dtypes)
```

Datatype of SkinThickness is: int64



We can see that there are 0 value data and SkinThickness cannot be 0. Hence replacing 0 with mean of SkinThickness class

```
[15]: data['SkinThickness']=data['SkinThickness'].replace(0,data['SkinThickness'].
↪mean())
```

```
[16]: data['Insulin'].values==0
```

```
[16]: array([ True,  True,  True, False, False,  True, False,  True, False,
        True,  True,  True,  True, False, False,  True, False,  True,
        False, False, False,  True,  True,  True, False, False,  True,
        False, False,  True,  True, False, False,  True,  True, False,
        True,  True,  True, False, False,  True,  True, False,  True,
        True,  True,  True,  True,  True, False, False, False, False,
        False,  True, False, False,  True, False,  True,  True,  True,
        False,  True,  True,  True,  True, False, False, False, False,
```



```

True, True, False, False, False, False, True, True, False,
True, True, False, False, False, True, True, False, True,
True, False, False, False, True, True, False, False, False,
True, True, True, True, False, False, False, True, False,
True, False, True, False, True, True, True, False, False,
False, False, True, False, False, False, False, False, False,
True, True, False, True, False, False, False, True, True,
True, True, True, False, False, False, True, False, False,
False, False, False, True, True, False, False, False, False,
False, True, True, True, True, True, True, True, False,
True, True, True, False, True, True, False, True, False,
False, False, True, False, True, False, True, True, True,
False, True, True, False, False, False, False, False, False,
False, True, False, True, True, False, True, True, False,
True, True, False, True, False, True, True, True, True,
True, False, True, False, True, True, True, False, False,
False, False, True, True, True, False, False, False, False,
False, True, False, False, False, True, False, False, False,
False, True, False, True, True, False, False, True, False,
True, True, False, False, False, True, False, False, True,
True, True, True, True, False, False, True, False, True,
True, False, True, True, False, False, True, True, False,
False, True, False, False, True, False, True, False, True,
True, False, True, False, False, True, True, False, True,
False, True, False, False, False, True, False, False, True,
False, False, True, True, False, True, False, True, False,
True, True, True, True, False, True, True, False, True,
False, True, True])

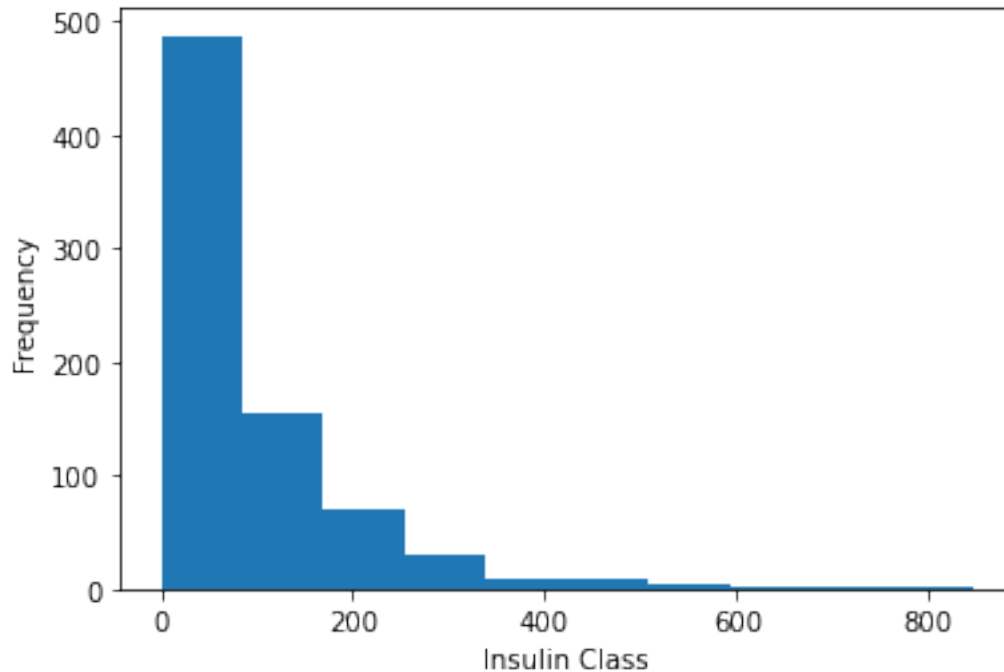
```

```

[17]: plt.xlabel('Insulin Class')
      data['Insulin'].plot.hist()
      print("Datatype of Insulin is:", data['Insulin'].dtypes)

```

Datatype of Insulin is: int64



We can see that there are 0 value data and Insulin cannot be 0.Hence replacing 0 with mean of Insulin class

```
[18]: data['Insulin']=data['Insulin'].replace(0,data['Insulin'].mean())
```

```
[19]: data['BMI'].values==0
```


[illegible]

[illegible]

```
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
True, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, True, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False]
```

```
[20]: plt.xlabel('BMI')
data['BMI'].plot.hist()
print("Datatype of BMI is:",data['BMI'].dtypes)
```

Datatype of BMI is: float64



We can see that there are 0 value data and BMI cannot be 0. Hence replacing 0 with mean of BMI class

```
[21]: data['BMI']=data['BMI'].replace(0,data['BMI'].mean())
```

```
[22]: data['Glucose'].value_counts().head(5)
```

```
[22]: 100.0    17
      99.0    17
      125.0   14
      106.0   14
      111.0   14
      Name: Glucose, dtype: int64
```

```
[23]: data['SkinThickness'].value_counts().head(5)
```

```
[23]: 20.536458    227
      32.000000     31
      30.000000     27
      27.000000     23
      23.000000     22
      Name: SkinThickness, dtype: int64
```

```
[24]: data['Insulin'].value_counts().head(5)
```

```
[24]: 79.799479     374
      105.000000     11
      130.000000      9
      140.000000      9
      120.000000      8
      Name: Insulin, dtype: int64
```

```
[25]: data['BloodPressure'].value_counts().head(5)
```

```
[25]: 70.0     57
      74.0     52
      68.0     45
      78.0     45
      72.0     44
      Name: BloodPressure, dtype: int64
```


```
[26]: data['BMI'].value_counts().head(5)
```

```
[26]: 32.000000     13
      31.600000     12
```



```
31.200000    12
31.992578    11
33.300000    10
Name: BMI, dtype: int64
```


```
[27]: data.dtypes.value_counts().plot(kind='bar')
plt.show()
```



We can see 6 float columns and 3 integer columns.

## WEEK 2: Data Exploration

```
[28]: (data.Outcome).value_counts().plot(kind = 'bar')
plt.show()
```




Data is Imbalanced

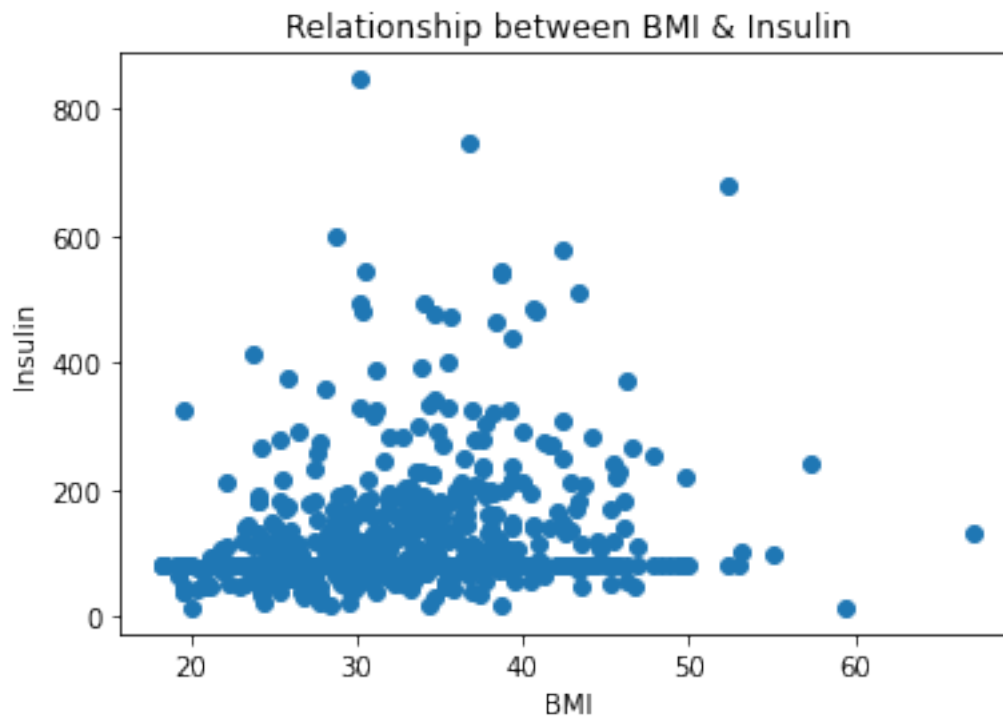
## 1 Scatter plot

```
[29]: BloodPressure = data['BloodPressure']  
      Glucose = data['Glucose']  
      SkinThickness = data['SkinThickness']  
      Insulin = data['Insulin']  
      BMI = data['BMI']
```

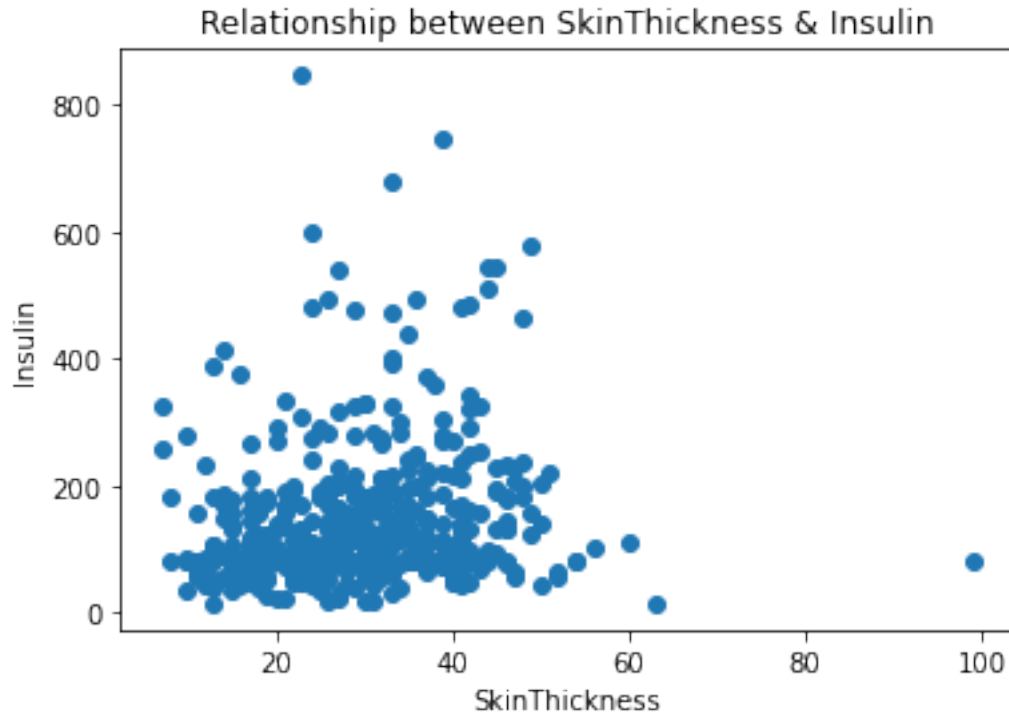
```
[30]: plt.scatter(BloodPressure, Glucose)  
      plt.xlabel('BloodPressure')  
      plt.ylabel('Glucose')  
      plt.title('Relationship between BloodPressure & Glucose')  
      plt.show()
```



```
[31]: plt.scatter(BMI, Insulin)
plt.xlabel('BMI')
plt.ylabel('Insulin')
plt.title('Relationship between BMI & Insulin')
plt.show()
```



```
[32]: plt.scatter(SkinThickness, Insulin)
plt.xlabel('SkinThickness')
plt.ylabel('Insulin')
plt.title('Relationship between SkinThickness & Insulin')
plt.show()
```



```
[33]: # correlation matrix
data.corr()
```

```
[33]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	\
Pregnancies	1.000000	0.127964	0.208984	0.013376	
Glucose	0.127964	1.000000	0.219666	0.160766	
BloodPressure	0.208984	0.219666	1.000000	0.134155	
SkinThickness	0.013376	0.160766	0.134155	1.000000	
Insulin	-0.018082	0.396597	0.010926	0.240361	
BMI	0.021546	0.231478	0.281231	0.535703	
DiabetesPedigreeFunction	-0.033523	0.137106	0.000371	0.154961	
Age	0.544341	0.266600	0.326740	0.026423	
Outcome	0.221898	0.492908	0.162986	0.175026	


  

	Insulin	BMI	DiabetesPedigreeFunction	\
Pregnancies	-0.018082	0.021546	-0.033523	
Glucose	0.396597	0.231478	0.137106	
BloodPressure	0.010926	0.281231	0.000371	
SkinThickness	0.240361	0.535703	0.154961	
Insulin	1.000000	0.189856	0.157806	
BMI	0.189856	1.000000	0.153508	
DiabetesPedigreeFunction	0.157806	0.153508	1.000000	
Age	0.038652	0.025748	0.033561	

Outcome	0.179185	0.312254	0.173844
	Age	Outcome	
Pregnancies	0.544341	0.221898	
Glucose	0.266600	0.492908	
BloodPressure	0.326740	0.162986	
SkinThickness	0.026423	0.175026	
Insulin	0.038652	0.179185	
BMI	0.025748	0.312254	
DiabetesPedigreeFunction	0.033561	0.173844	
Age	1.000000	0.238356	
Outcome	0.238356	1.000000	

```
[34]: sns.heatmap(data.corr(),annot=True)
```

```
[34]: <AxesSubplot:>
```



```
[35]: data.corr().style.background_gradient()
```

```
[35]: <pandas.io.formats.style.Styler at 0x17064a4d550>
```

### WEEK 3: Data Modeling

```
[36]: data.head()
```

```
[36]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148.0	72.0	35.000000	79.799479	33.6	
1	1	85.0	66.0	29.000000	79.799479	26.6	
2	8	183.0	64.0	20.536458	79.799479	23.3	
3	1	89.0	66.0	23.000000	94.000000	28.1	
4	0	137.0	40.0	35.000000	168.000000	43.1	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
[37]: x=data.iloc[:,[0,1,2,3,4,5,6,7]].values
      y=data['Outcome'].values
```

```
[38]: from sklearn.model_selection import train_test_split
```

```
[39]: x_train,x_test,y_train,y_test=train_test_split(x,y,random_state=0,test_size=0.2)
```

### LogisticRegression

```
[40]: from sklearn.linear_model import LogisticRegression
      model = LogisticRegression()
      model.fit(x_train,y_train)
```

G:\Software\anaconda\lib\site-packages\sklearn\linear\_model\\_logistic.py:762:  
ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
[40]: LogisticRegression()
```

```
[41]: print(model.score(x_train,y_train))
      print(model.score(x_test,y_test))
```

0.758957654723127  
0.8311688311688312

```
[42]: from sklearn.metrics import classification_report
print(classification_report(y,model.predict(x)))
```

	precision	recall	f1-score	support
0	0.80	0.87	0.83	500
1	0.71	0.59	0.64	268
accuracy			0.77	768
macro avg	0.76	0.73	0.74	768
weighted avg	0.77	0.77	0.77	768

#### DecisionTreeClassifier

```
[43]: from sklearn.tree import DecisionTreeClassifier
model2 = DecisionTreeClassifier(max_depth=5)
model2.fit(x_train,y_train)
```

```
[43]: DecisionTreeClassifier(max_depth=5)
```

```
[44]: print(model2.score(x_train,y_train))
print(model2.score(x_test,y_test))
```

0.8208469055374593  
0.7662337662337663

```
[45]: from sklearn.metrics import classification_report
print(classification_report(y,model2.predict(x)))
```

	precision	recall	f1-score	support
0	0.89	0.81	0.85	500
1	0.69	0.81	0.75	268
accuracy			0.81	768
macro avg	0.79	0.81	0.80	768
weighted avg	0.82	0.81	0.81	768

#### KNN

```
[46]: from sklearn.neighbors import KNeighborsClassifier
model3 = KNeighborsClassifier(n_neighbors=7,
                             metric='minkowski',
                             p = 2)
```



```
model3.fit(x_train,y_train)
```

```
[46]: KNeighborsClassifier(n_neighbors=7)
```

```
[47]: print(model3.score(x_train,y_train))  
print(model3.score(x_test,y_test))
```

```
0.7899022801302932  
0.7337662337662337
```

```
[48]: from sklearn.metrics import classification_report  
print(classification_report(y,model3.predict(x)))
```

	precision	recall	f1-score	support
0	0.81	0.86	0.83	500
1	0.71	0.63	0.66	268
accuracy			0.78	768
macro avg	0.76	0.74	0.75	768
weighted avg	0.77	0.78	0.78	768

#### WEEK 4: Data Modeling

```
[49]: # ROC Curve KNN  
from sklearn.metrics import roc_curve  
from sklearn.metrics import roc_auc_score  
  
probs = model3.predict_proba(x)  
# keep probabilities for the positive outcome only  
probs = probs[:, 1]  
# calculate AUC  
auc = roc_auc_score(y, probs)  
print('AUC: %.3f' % auc)  
# calculate roc curve  
fpr, tpr, thresholds = roc_curve(y, probs)  
print("True Positive Rate - {}, False Positive Rate - {} Thresholds - {}".  
      ↪format(tpr,fpr,thresholds))  
# plot the roc curve for the model  
plt.plot([0, 1], [0, 1], linestyle='--')  
plt.plot(fpr, tpr, marker='.')  
plt.xlabel("False Positive Rate")  
plt.ylabel("True Positive Rate")
```

```
AUC: 0.845
```


```
True Positive Rate - [0.          0.05970149 0.25746269 0.45149254 0.62686567  
0.80970149
```

```

0.91791045 0.99253731 1.          ], False Positive Rate - [0.      0.004 0.014
0.064 0.14  0.27  0.474 0.7    1.    ] Thresholds - [2.          1.
0.85714286 0.71428571 0.57142857 0.42857143
0.28571429 0.14285714 0.          ]

```

[49]: Text(0, 0.5, 'True Positive Rate')



```

[50]: #Precision Recall Curve for Logistic Regression

from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score
# predict probabilities
probs = model.predict_proba(x)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# predict class values
yhat = model.predict(x)
# calculate precision-recall curve
precision, recall, thresholds = precision_recall_curve(y, probs)
# calculate F1 score
f1 = f1_score(y, yhat)
# calculate precision-recall AUC
auc = auc(recall, precision)

```


```

# calculate average precision score
ap = average_precision_score(y, probs)
print('f1=%.3f auc=%.3f ap=%.3f' % (f1, auc, ap))
# plot no skill
plt.plot([0, 1], [0.5, 0.5], linestyle='--')
# plot the precision-recall curve for the model
plt.plot(recall, precision, marker='.')

```

f1=0.643 auc=0.727 ap=0.728

[50]: [<matplotlib.lines.Line2D at 0x170652f4820>]



```

[51]: #Precision Recall Curve for KNN

from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score
# predict probabilities
probs = model3.predict_proba(x)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# predict class values
yhat = model3.predict(x)
# calculate precision-recall curve
precision, recall, thresholds = precision_recall_curve(y, probs)

```


```

# calculate F1 score
f1 = f1_score(y, yhat)
# calculate precision-recall AUC
auc = auc(recall, precision)
# calculate average precision score
ap = average_precision_score(y, probs)
print('f1=%.3f auc=%.3f ap=%.3f' % (f1, auc, ap))
# plot no skill
plt.plot([0, 1], [0.5, 0.5], linestyle='--')
# plot the precision-recall curve for the model
plt.plot(recall, precision, marker='.')

```

f1=0.664 auc=0.750 ap=0.713

[51]: [<matplotlib.lines.Line2D at 0x17065350a90>]



[52]: #Precision Recall Curve for Decision Tree Classifier

```

from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score
# predict probabilities
probs = model2.predict_proba(x)
# keep probabilities for the positive outcome only
probs = probs[:, 1]

```


```

# predict class values
yhat = model2.predict(x)
# calculate precision-recall curve
precision, recall, thresholds = precision_recall_curve(y, probs)
# calculate F1 score
f1 = f1_score(y, yhat)
# calculate precision-recall AUC
auc = auc(recall, precision)
# calculate average precision score
ap = average_precision_score(y, probs)
print('f1=%.3f auc=%.3f ap=%.3f' % (f1, auc, ap))
# plot no skill
plt.plot([0, 1], [0.5, 0.5], linestyle='--')
# plot the precision-recall curve for the model
plt.plot(recall, precision, marker='.')

```

f1=0.749 auc=0.772 ap=0.779

[52]: [<matplotlib.lines.Line2D at 0x170653aad30>]



[ ]: